

Rao Vinnakota  
Sven Anderson  
Machine Learning  
HW # 3

I divided the problem into two portions - pruning the node and choosing when to prune.

### **Pruning the Node:**

I wrote the code to prune the node without using pop(). I thought about what it meant to prune a node:

1. The node is now a leaf
2. The value of the leaf is now the major classification.

To prune the node, change the value of tree['@LEAF'] to True. Then create tree['@VALUE'] to be equal to tree['@MAJOR']. Now, to make sure that the node becomes a leaf, the algorithm iterates through the features and deletes them. Last, if it needs to, it deletes the attribute

### **Choosing Nodes to Prune:**

I used the algorithm of reduced-error pruning as the basis for how I pruned the mush decision tree. The algorithm worked in these broad steps:

1. Create a new test tree
2. Prune the current node of the test tree
3. If that pruned test tree scores higher than the original tree
  - a. prune the original tree and run the function with the new pruned tree
4. If the pruned test tree doesn't score higher:
  - a. Iterate through the features of the current node calling the prune function for each node that it leads to.

For the prune\_mush to work, I copied score\_tree into another function return\_score. In return score, the percentage correct is returned instead of printed out.

### **Results:**

Pruning in this case was very effective. The percentage correct on the validation set rose from 80 to 86% while going from 80 to 89% on the testing set. An unwieldy tree with 1153 nodes was also lowered to a far more manageable 51 nodes. It may be better to traverse the tree a different way, rather than top-down, for a more effective attempt.